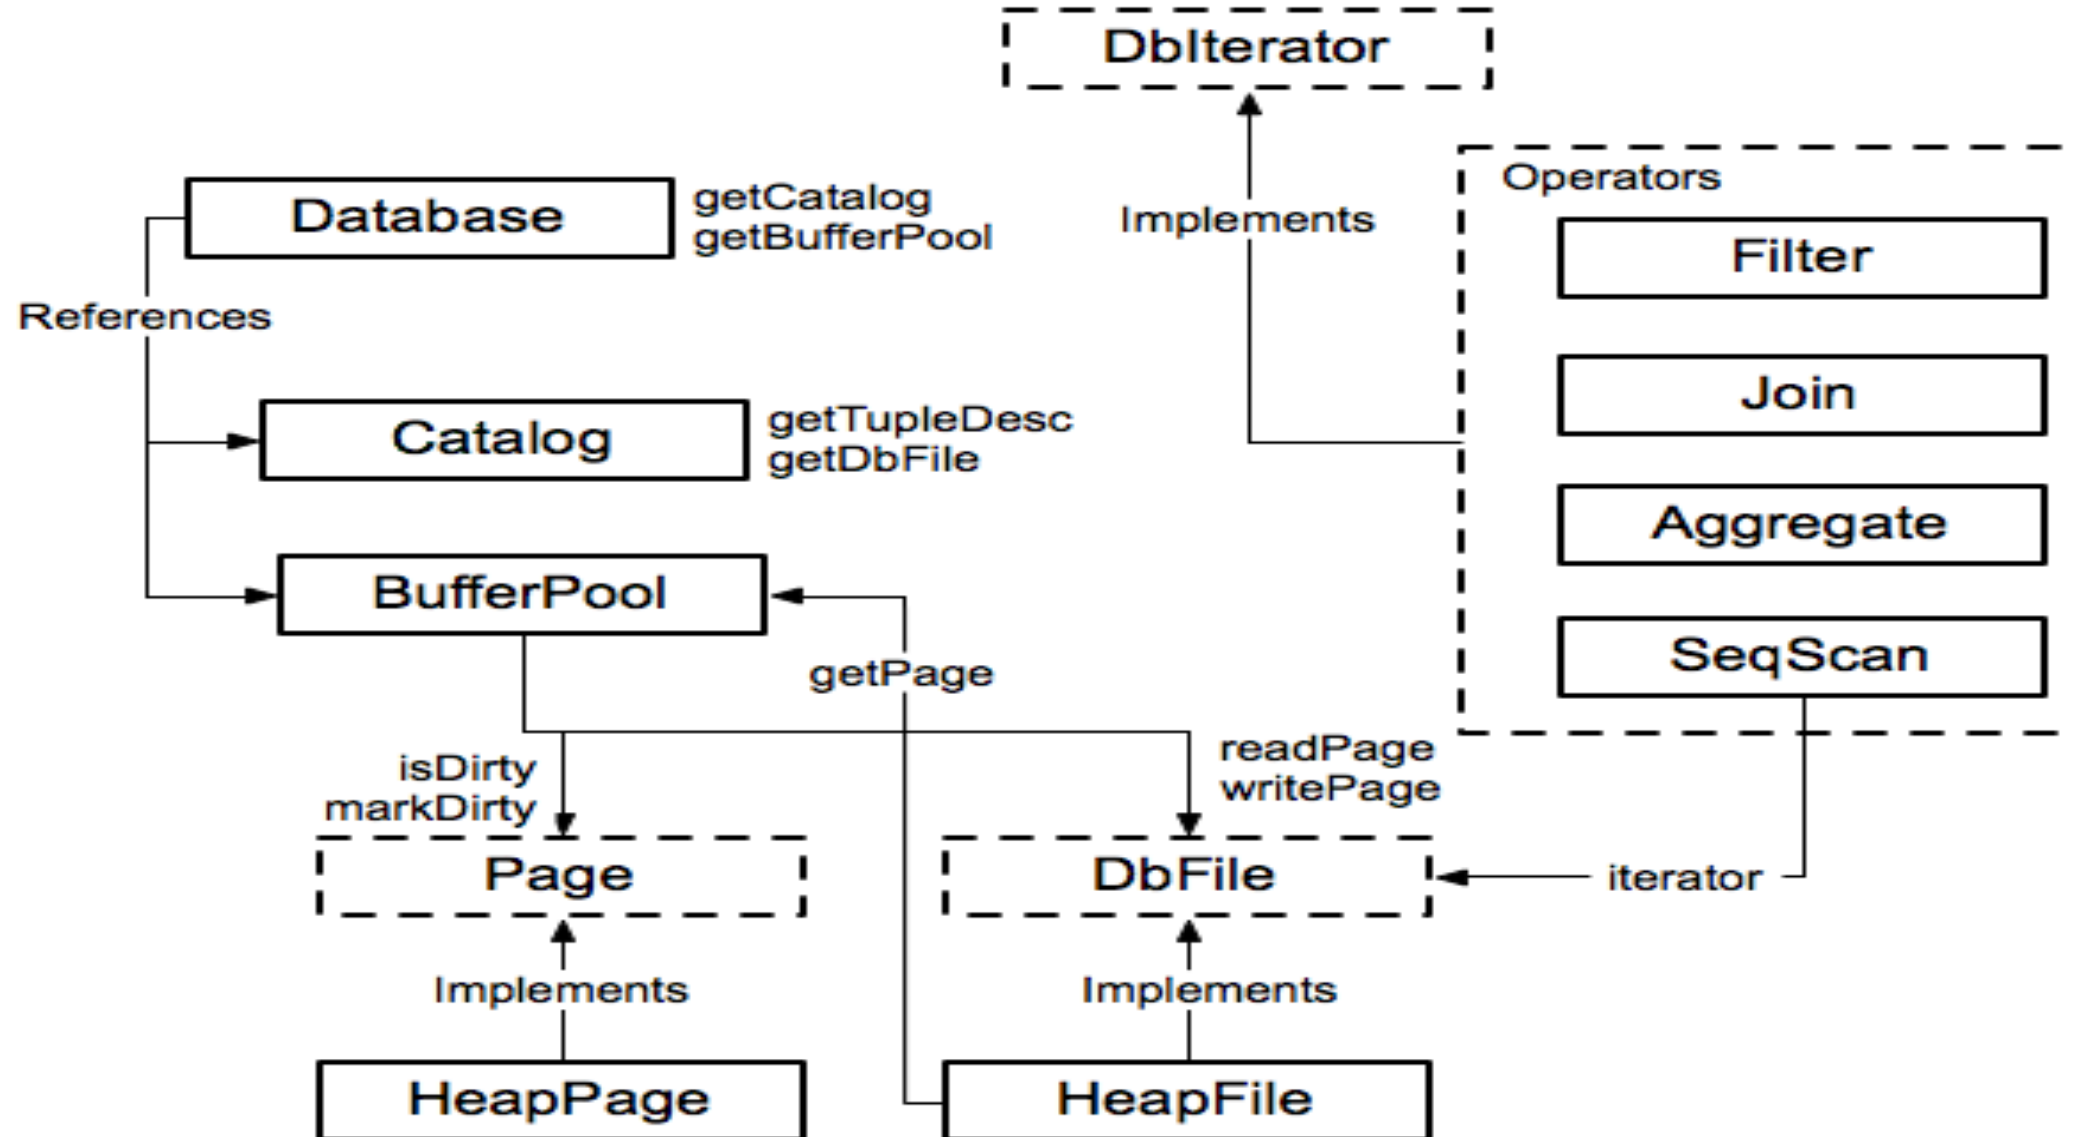


# What is SimpleDB?

---

- A basic database system developed by Sam Madden and others from MIT
- SQL Front-end (Provided for later labs)
  - Heap files (Lab 1)
  - Buffer Pool (Labs 1-6)
  - Basic Operators (Labs 1 & 2)
    - Scan, Filter, JOIN, Aggregate
  - B-Tree Indexes (Lab 3)
  - Transactions (Lab 4)
  - Query optimizer (Lab 5)
  - Recovery (Lab 6)
- Javadoc is your friend!

# Module Diagram



# Database

---

## Singleton Database:

**Database.getCatalog()**

**Catalog**  
**=> List of DB tables**

**Database.getBufferPool()**

**BufferPool**  
**=> Caches DB pages in memory**

**Database.getLogFile()**

**LogFile**  
**(Ignore for Lab 1)**

# Catalog

---

**Catalog:**

<b>DbFile ID</b>	<b>Table</b>
001	Table1
002	Table2
003	Table3

**Table:**

**DbFile file**  
**String name**  
**String primary\_key**

**=> Stores a list of all tables in the database**

# BufferPool

**Buffer Pool:**

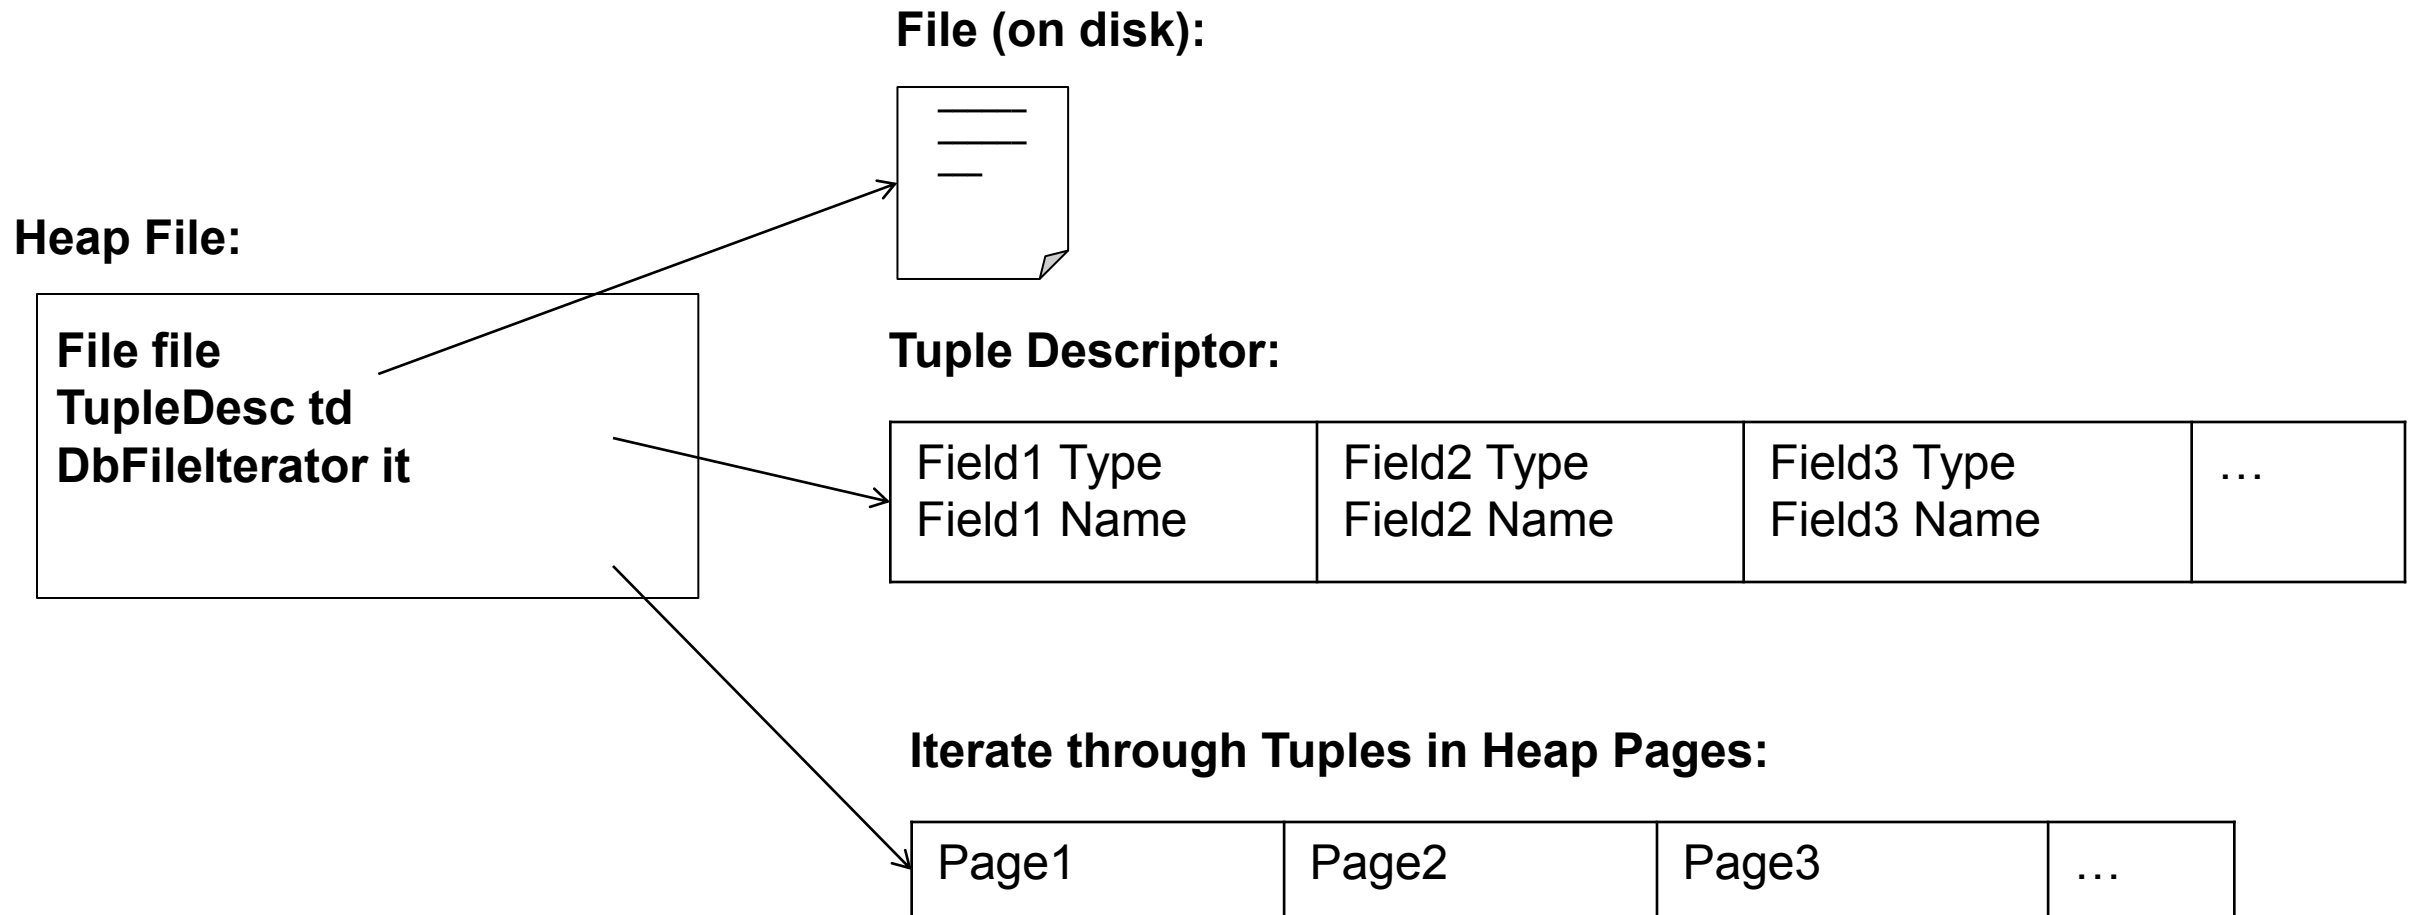
Page ID	Page
001	Page1
003	Page3
007	Page7

**Page:**

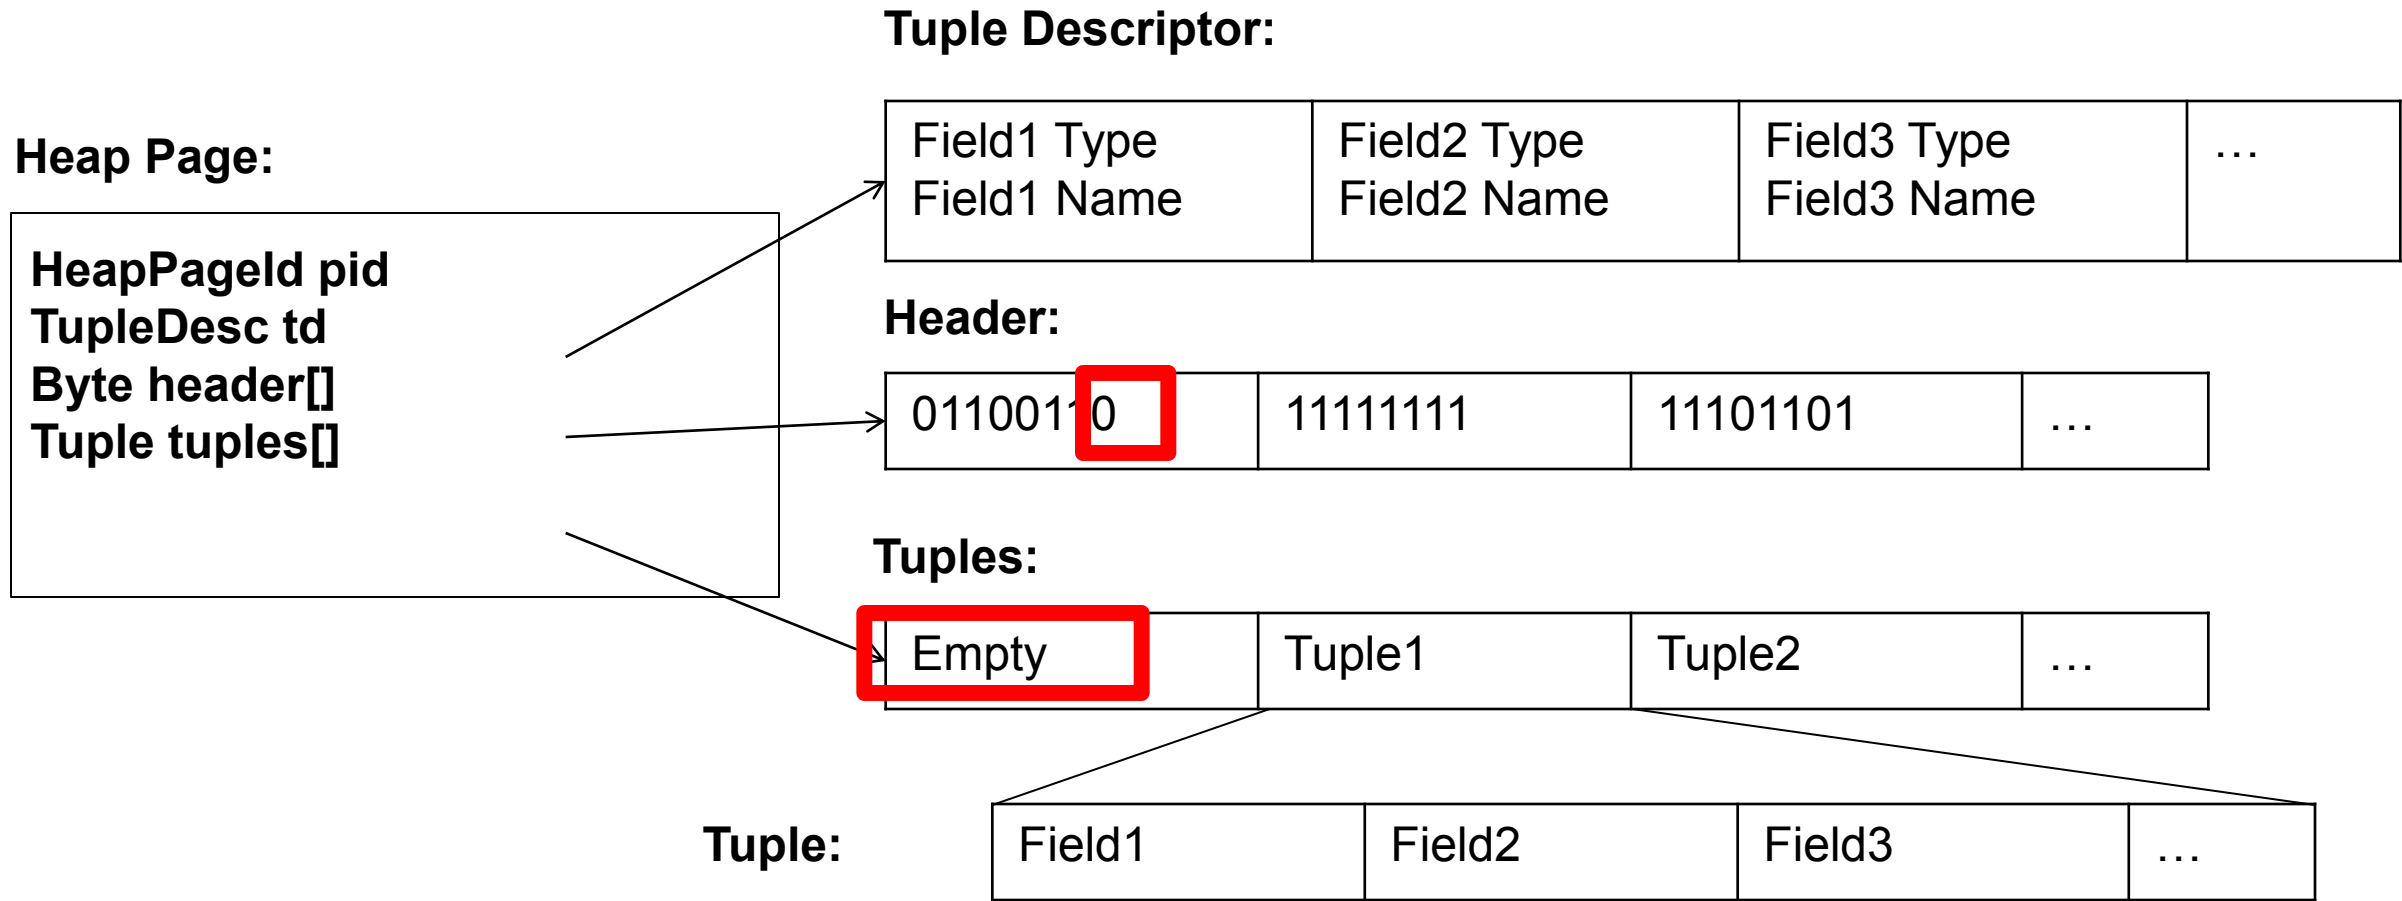
**Pageld id**  
**Tuple tuples[]**  
**Byte header[]**

**=> Caches recently accessed database pages in memory**

# HeapFile (Implements DbFile)



# HeapPage (Implements Page)



***Fields and Tuples are Fixed Width!***

# SeqScan (Implements DbIterator)

---

- DbIterator class implemented by all operators
  - open()
  - close()
  - getTupleDesc()
  - hasNext()
  - next()
  - rewind()
- Iterator model: chain iterators together
  - Use DbFileIterator from HeapFile



```
// construct a 3-column table schema
```

```
Type types[] = new Type[]{ Type.INT_TYPE, Type.INT_TYPE, Type.INT_TYPE };
```

```
String names[] = new String[]{ "field0", "field1", "field2" };
```

```
TupleDesc descriptor = new TupleDesc(types, names);
```

```
// create the table, associate it with some_data_file.dat
```

```
// and tell the catalog about the schema of this table.
```

```
HeapFile table1 = new HeapFile(new File("some_data_file.dat"), descriptor);
```

```
Database.getCatalog().addTable(table1);
```

```
// construct the query: we use a simple SeqScan, which spoonfeeds
```

```
// tuples via its iterator.
```

```
TransactionId tid = new TransactionId();
```

```
SeqScan f = new SeqScan(tid, table1.id());
```

```
// and run it
```

```
f.open();
```

```
while (f.hasNext()) {
```

```
    Tuple tup = f.next();
```

```
    System.out.println(tup);
```

```
}
```

```
f.close();
```

```
Database.getBufferPool().transactionComplete();
```

# HeapFileEncoder.java

---

- Because you haven't implemented insertTuple, you have no way to create data files
- HeapFileEncoder converts CSV files to HeapFiles
- Usage:
  - `java -jar dist/simplydb.jar convert csv-file.txt numFields`
- Produces a file `csv-file.dat`, that can be passed to HeapFile constructor.

# Compiling, Testing, and Running

---

- Demo on running tests and debugging in Ubuntu and with Eclipse